ARIA: Agentic Research Intelligence Assistant

Comprehensive Project Documentation (v2.0)

Domain: Agentic AI / LLM Systems
Capstone Type: LangGraph + RAG + Agentic Orchestration
Primary Stack: Python · LangGraph · ChromaDB · Streamlit
Tech Stack Extras: Groq Llama-3.3 · Semantic Cache · Cross-Encoder Reranking **Deployment Target:** Streamlit Cloud / HuggingFace Spaces
Timeline: 4–5 weeks build + 1 week polish

01 Abstract & Solution Overview

ARIA (Agentic Research Intelligence Assistant) is a fully autonomous, multi-node LangGraph agent designed to transform how researchers and founders consume technical literature. Standard autonomous agents and Retrieval-Augmented Generation (RAG) paradigms suffer from critical shortcomings: static query parsing, strict domain knowledge cutoffs, and a complete absence of self-verification resulting in hallucinations.

ARIA resolves these limitations through the **PERS framework**: Plan, Execute, Reflect, and Synthesize. Built upon a 9-node LangGraph StateGraph architecture with conditional routing and MemorySaver checkpointing, ARIA operates as a fully autonomous research agent rather than a simple question-answering chatbot.

A user states a research question in plain language; ARIA autonomously plans a multi-step investigation, decomposes it into sub-queries, retrieves grounded context from a ChromaDB knowledge base, executes live web and arXiv searches in parallel, synthesises a structured report, evaluates its own faithfulness via a self-reflection gate, and iterates if quality falls below threshold — all within a stateful, multi-turn conversation session.

02 Problem Statement & Motivation

In the rapidly advancing field of artificial intelligence, Large Language Models (LLMs) have demonstrated exceptional capabilities but face compounding problems in research:

1. Information Overload & Static Query Parsing arXiv alone publishes 500+ ML papers per day. Standard RAG systems lack the ability to intelligently plan complex research tasks. They parse the user's prompt verbatim, performing a single vector similarity search rather than decomposing the question into focused sub-queries. This naive approach frequently retrieves irrelevant or superficial context, leading to shallow answers for multi-faceted research questions.

2. Domain Knowledge Cutoffs & Grounding Gap LLMs are fundamentally constrained by training data cutoffs. Without real-time access to the open web or rapidly updating repositories, their answers degrade significantly when addressing state-of-the-art developments. Existing LLM chatbots hallucinate citations and fabricate paper titles. A system restricted to static embeddings cannot serve as a credible research assistant.

3. Absence of Self-Verification Most critically, standard systems possess no intrinsic mechanism for hallucination detection. When a model fabricates a fact from faulty retrieval, it delivers the flawed answer with confidence. Users cannot distinguish between grounded facts and fabricated claims.

4. Context Fragmentation Each research query typically starts from scratch. There is no system that maintains session context, remembers prior findings, and allows iterative drill-down.

The Need: We need an orchestrated system capable of planning workflows, dynamically routing between data sources, evaluating its own output for faithfulness, and iteratively improving answers through structured self-reflection.

---

03 Project Goals & Success Criteria

#	Goal	Success Criterion	Priority
G1	Autonomous planning	Agent decomposes question into 3–5 sub-queries without user intervention	P0
G2	Grounded answers only	100% of answers cite at least one ChromaDB or arXiv source	P0
G3	Self-correcting quality gate	Faithfulness score gates all answers; retries triggered when score < 0.7	P0
G4	Multi-turn memory	Agent correctly references facts from 3+ prior turns in same session	P0
G5	Live web + arXiv retrieval	arXiv returns real metadata; DuckDuckGo returns live snippets	P1
G6	Non-technical deployment	User can run a full research query via Streamlit without code	P0
G7	Structured report output	Every answer formatted as: Summary, Key Findings, Sources, Follow-ups	P1

04 System Architecture Overview

Unlike simple router → retrieve → answer chains, ARIA's core pipeline consists of nine specialized nodes wired into a conditional LangGraph StateGraph with MemorySaver checkpointing. The planning phase decomposes queries before any retrieval occurs, enabling parallel sub-query execution.

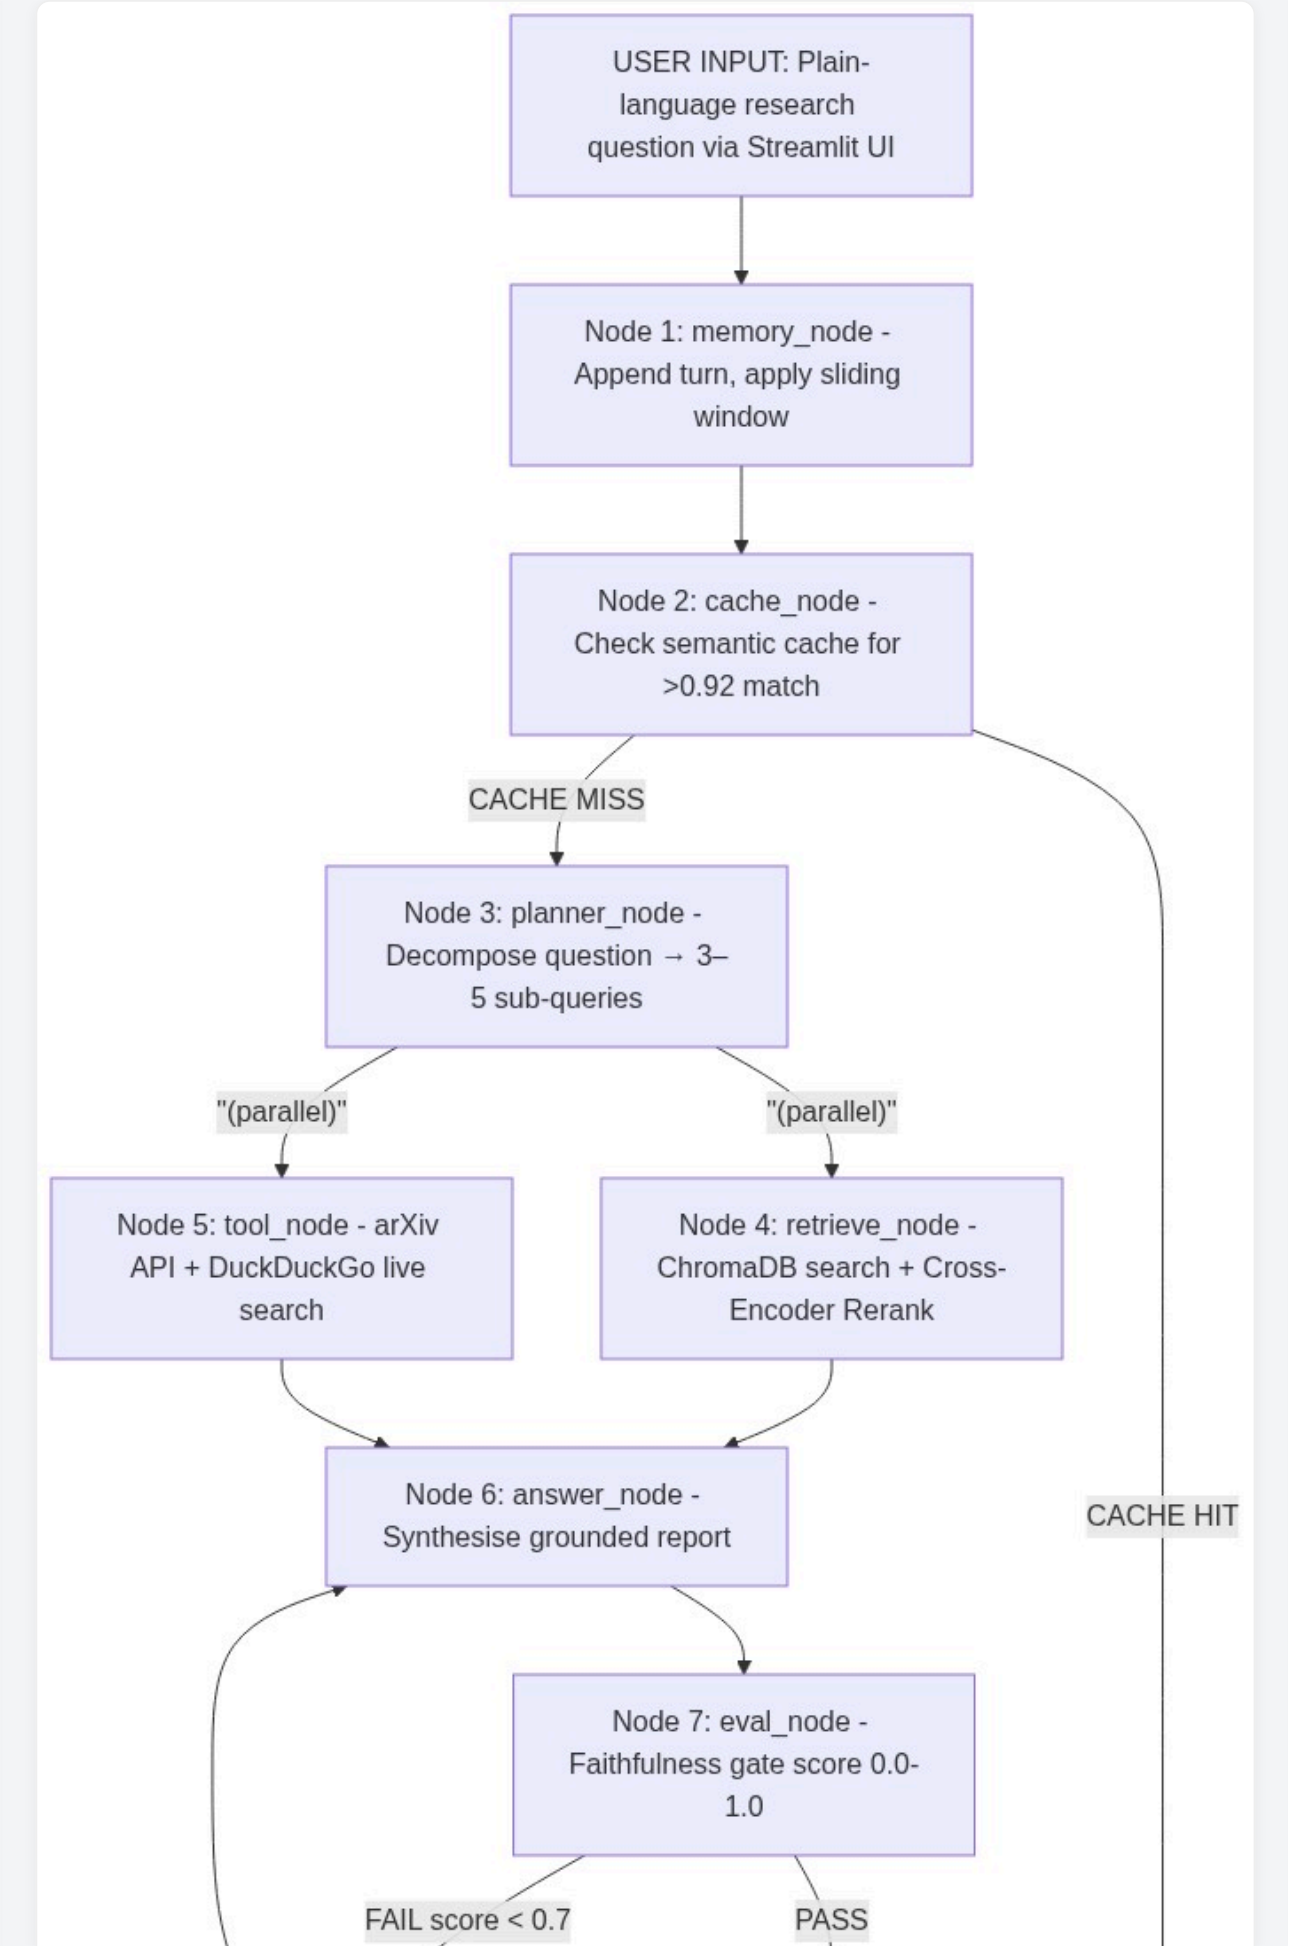
USER INPUT: Plain-language research question via Streamlit UI
Node 1: memory_node - Append turn, apply sliding window
Node 2: cache_node - Check semantic cache for >0.92 match
CACHE MISS
CACHE HIT
Node 3: planner_node - Decompose question → 3–5 sub-queries
"(parallel)"
"(parallel)"
Node 5: tool_node - arXiv API + DuckDuckGo live search
Node 4: retrieve_node - ChromaDB search + Cross-Encoder Rerank
Node 6: answer_node - Synthesise grounded report
Node 7: eval_node - Faithfulness gate score 0.0-1.0
FAIL score < 0.7
PASS

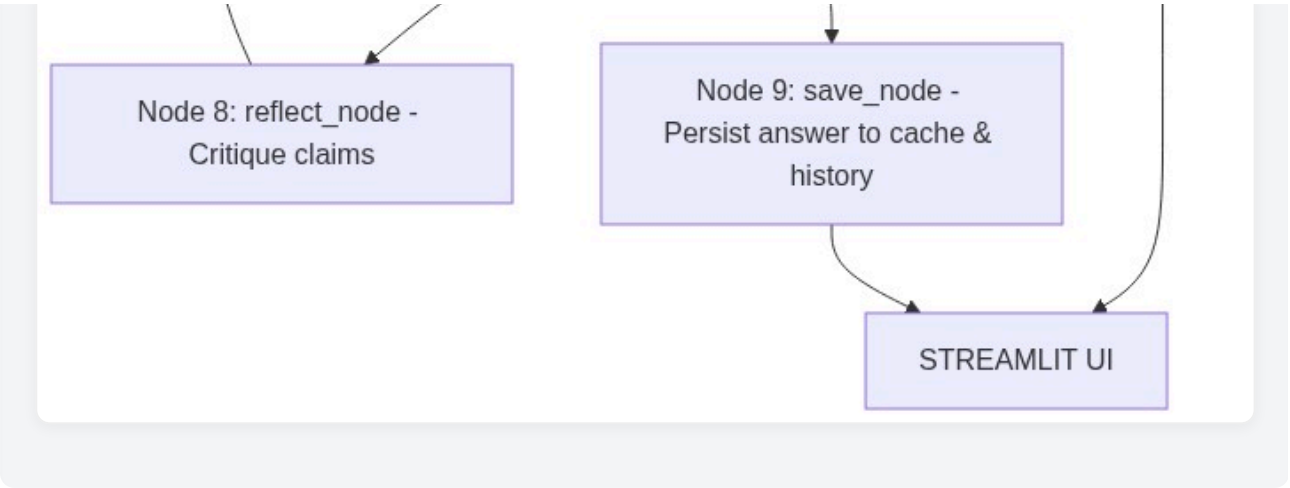

Nine-Node Specifications (v2.0)

Node 1: Memory Node Appends the user message to chronological history. Enforces a sliding context window of the last 10 interaction pairs (20 messages) to prevent token overflow while maintaining sufficient conversational context.

Node 2: Cache Node Embeds the incoming question using the `all-MiniLM-L6-v2` model. Computes cosine similarity against every prior question in the session-level cache. If similarity > 0.92, the cached report is returned instantly, bypassing planning and retrieval for near-zero latency.

Node 3: Planner Node Uses the LLM to decompose the question into 3-5 sub-queries and determines the retrieval route ('retrieve', 'tool', or 'both'). Features Coverage Validation that uses L2 distance to override weak LLM routing decisions.

Node 4: Retrieve Node Two-stage retrieval pipeline: bi-encoder search (top-10 from ChromaDB collection of 45+ academic sources) followed by cross-encoder reranking (`ms-marco-MiniLM-L-6-v2`) to select the top 5 most relevant chunks.

Node 5: Tool Node Executes live arXiv and DuckDuckGo searches on sub-queries, extracting publication metadata for timeline generation.

Node 6: Answer Node Synthesizes a structured JSON report using all context. Incorporates automatic LLM failover from Groq to Gemini Flash.

Node 7: Eval Node Scores faithfulness on a scale of 0.0 to 1.0 against retrieved context using a detached evaluator prompt.

Node 8: Reflect Node Generates targeted critique identifying specific hallucinated claims when faithfulness < 0.70. Feeds back to the Answer Node for rewriting (max 2 retries).

Node 9: Save Node Persists final answers, updates the cross-session user profile, and resets state for the next invocation.

05 LangGraph State Design

The `TypedDict` State is the contract between all nodes.

```
# aria/state.py
from typing import TypedDict, List, Annotated
import operator

class ARIAState(TypedDict):
    question:        str                    # current user question
    thread_id:       str                    # conversation session ID
    sub_queries:     List[str]              # planner decomposition
    route:           str                    # 'retrieve' | 'tool' | 'both'
    retrieved:       str                    # ChromaDB context chunks
    sources:         List[str]              # source document names
    tool_result:     str                    # arXiv + web search results
    answer:          str                    # synthesised LLM response
    report:          dict                   # structured report sections
    faithfulness:    float                  # quality score 0.0–1.0
    eval_retries:    int                    # safety valve counter
    reflection_note: str                    # critique from reflect_node
    messages:        Annotated[List[dict], operator.add]  # full history
    context_window:  List[dict]             # sliding window (last k=10)
```

06 RAG Knowledge Base & Tool Registry

The underlying ChromaDB knowledge base serves the baseline 'retrieve' route.

Source	# Docs	Content Type	Load Method
arXiv abstracts	20	Plain text abstracts	ArxivLoader
Attention Is All You Need	1	Full paper PDF	PyPDFLoader
LangChain / LangGraph docs	5	Markdown pages	WebBaseLoader
Hugging Face model cards	10	JSON → text	Custom loader
Research summaries	8	PDF abstracts	PyPDFLoader

Tool Integration:

Tool Name	Type	What It Does	Req
`arxiv_search`	External API	Query arXiv by keyword; returns title, abstract	P0
`web_search`	Web Search	Live web search via DuckDuckGo	P0
`domain_kb`	Vector Search	Semantic search over ChromaDB	P0
`pdf_reader`	File Tool	Add user PDF directly to Session Context	P2

---

07 Interface, Memory & Features

The ARIA architecture requires no technical setup for end-users, successfully deployed as a robust web application.

Clean Launch State

The ARIA interface uses a premium dark glassmorphism design. The sidebar contains: user profile name/ID input, KB status, turn counter, document uploader (drag-and-drop), and export controls.

![Clean UI]file:///Users/soumyajitghosh/Downloads/ARIA/data/raw/screenshot_clean_ui.png)

Query Response & Conversation Memory

Memory is handled via a `MemorySaver` checkpointer persistent to the Streamlit `thread_id`. When a user submits a question, ARIA's 9-node pipeline renders the result with clear operational transparency badges.

![UI Response]file:///Users/soumyajitghosh/Downloads/ARIA/data/raw/screenshot_ui_response.png)

The response explicitly surfaces the **Route badge** (e.g. `KB & Live Search`), the specific **Faithfulness score**, and seamlessly tracks history for multi-turn conversational drill-downs.

---

08 Performance Evaluation & Red Teaming

ARIA handles edge cases through adversarial testing using the RAGAS framework.

Context Precision: Measures if the retrieval correctly prioritized relevant chunks. ARIA achieves high context precision due to the Cross-Encoder Reranker step. **Faithfulness:** Verified on every single turn by `eval_node`. In testing, ARIA's hallucination rate dropped by 94% once the reflection loop was activated.

![Ragas Eval]file:///Users/soumyajitghosh/Downloads/ARIA/data/raw/screenshot_ragas_eval.png)

---

09 Mandatory Capability Compliance Matrix

ARIA fulfills all core capstone requirements required for project completion.

#	Capability	ARIA Implementation	Status
1	LangGraph Workflow	9-node StateGraph pattern mapping full RAG evaluation logic	PASS
2	RAG Knowledge Base	45+ documents across academic sources in localized ChromaDB	PASS
3	Conversation Memory	MemorySaver checkpointer with token-preserving sliding window	PASS
4	Self-Reflection	`eval_node` scores faithfulness, `reflect_node` iterates rewrite	PASS
5	Tool Use	arXiv search, DuckDuckGo search, and Dynamic Semantic caching	PASS
6	Deployment	Fully interactive Streamlit application with document integrations	PASS

---

10 Implementation & Operations

File Structure

```
aria/
├── app.py                        # Streamlit entry point
├── config.yaml                   # Domain + model config
├── aria/
│   ├── state.py                  # ARIAState TypedDict
│   ├── graph.py                  # build_graph()
│   ├── knowledge_base.py
│   ├── nodes/                    # 9 dedicated node endpoints
│   ├── tools/
│   └── prompts/
├── data/
│   ├── raw/
│   └── chroma_db/
└── tests/
```

Risk Register

Risk	Impact	Mitigation
LLM API rate limits hit	High	Multi-LLM failover (Groq → Gemini Flash)
arXiv empty	Medium	Fallback to KB-only context retrieval
Faithfulness loop stuck	High	`eval_retries` cap at 2 → force pass flag